

Distributed Flow Control and Intelligent Data Transfer in High Performance Computing Networks

Hochschule Offenburg

Mehdi Sadeghi

13. Feb 2015

Declaration

I hereby declare and confirm that this thesis is entirely the result of my own original work. Where other sources of information have been used, they have been indicated as such and properly acknowledged. I further declare that this or similar work has not been submitted for credit elsewhere.

Offenburg, February 28, 2015
Mehdi Sadeghi

Abstract

This document contains my master's thesis report, including the problem definition, an overview of state of the art, discussions and my suggestions.

Acknowledgement

Here will come acknowledgement

Contents

1	Introduction	7
1.1	Objectives	7
1.2	Terminology	7
1.3	Problem Context	8
1.4	Assumptions	9
1.4.1	Collaborating Network	9
1.4.2	Data Characteristics	9
1.4.3	Data Transfer	9
1.4.4	Workflow	9
2	Rough Ideas	11
2.1	Intelligent Data Transfer: Use Case One	11
2.1.1	Identical Instances	11
2.1.2	The Idea	12
3	Literature	13
3.1	Parameters	13
3.2	Data Storage Systems	14
3.3	Distributed File Systems	14
3.3.1	Hadoop Distributed File System (HDFS)	14
3.3.2	XTREEMFS	15
4	Scenarios	16
4.1	Operation	16
4.1.1	Linear Operation	16
4.1.2	Non-linear Operation	16
4.2	Decision Making	17
4.2.1	Scenario 1	18
5	Workflow Management	19

6 Data Transfer	20
References	21

Preface

European scientific communities launch many scientific experiments daily, resulting in huge amounts of data. Specifically in molecular dynamics and material science fields there are many different simulation software which are being used to accomplish multiscale modelling tasks. These tasks often involve running multiple simulation programs over the existing data or the data which is produced by other simulation software. Depending on the amount of the data and the desired type of simulation these tasks could take many days to finish. The order to run simulation software and providing the input data are normally defined in scripts written by researchers which is the simplest form of workflow management.

While small experiments could be handled with simple scripts and normal computers, larger scale experiments require different solutions. These type of experiments demand huge computing power which are made available by computer clusters and super computers. An important characteristic of larger experiments is the amount of produced data. This data, which needs to be transferred many times back and forth between computers, grow by an order of magnitude, resulting in terabytes of data.

Large scientific experiments are the source of many high performance computing (HPC) problems, specially data transfer and workflow management. Moreover HPC resources are expensive and should be used efficiently, therefore making data transfer more efficient is important.

This thesis is an effort to know the main data transfer scenarios in HPC experiments and try to address them in a distributed manner with a collective but decentralized approach toward workflow management.

If there are any comments and improvements regarding this document, the author appreciates an email to the following address:

`msadeghi@stud.hs-offenburg.de`
Hochschule Offenburg
Mehdi Sadeghi

Chapter 1

Introduction

1.1 Objectives

We will focus on two important topics in this work. First one is data transfer problems between multiple computers which are doing a task collaboratively. Second one is the collaboration itself, i.e. how multiple computers will accomplish a collaborative task in a distributed and decentralized environment. To accomplish these objectives we will discuss our specific distributed work flow management and data transfer scenarios. We define our requirements regarding the above mentioned topics and we will extract the parameters which we are going to assess other solutions with them. Then we will go through the currently available solutions. Afterward we will discuss their applicability according to our requirements and whether they answer our needs or not. The goal of this thesis is to minimize the amount of transferred data in a network of collaborating computers. This data belongs to operations which might be linear or not and require multiple steps. The operation can be initiated in any participating computer but the required data is not necessarily available on that computer, even though the operation result should be delivered back to the reinitializing computer.

1.2 Terminology

We will use a number of terms through this report. Here are the meaning for each.

Node Each node refers to one computer in the network.

Data When we refer to data we mean the output of scientific applications, such as NumPy types.

Dataset Same as data with more emphasize on it as collection of e.g. NumPy types.

Application Refers to the demo application which has been developed to show case the proposed solution.

Instance Refers to an instance of the same application running on a node.

Operation Any sort of operation, linear or non-linear which is being provided by the application.

Task Same as the operation with more emphasize on the output rather than the functionality.

Service A scientific operation being provided by the application which could be called remotely.

System The combination of nodes, data, application, instances, operations and services as a whole.

User A scientist, researcher or student who uses the system to run a task.

1.3 Problem Context

Whole this work is an effort to address issues of a scientific environment. Some particular characteristics are running multiple scientific programs on different computers which need to exchange data in order to accomplish one operation. Another task which is often done is visualization. Visualizing the operation results ,depending on the requested visualization, might require heavy computational tasks i.e. average or comparison on data which might not be available on the same machine or might be residing partially on different computers. The produced data often exceeds 1 GB in many experiments and it should be moved back and forth every few minutes, therefore it is cheaper to transfer the operation rather than the data.

The problem here is not about distributing the stored data rather, data exchange between instances of the application talking together in runtime while doing one global task and keeping this workflow distributed. In this terms each application instance takes care of its own data and provides a set of services. Some operations require data from another node, therefore we have to transfer the data or run the operation on the node which contains the data. There are a number of scenarios which we will discuss.

1.4 Assumptions

During this work we have a number of assumptions. We have a certain problem which we want to focus on rather than reintroducing solutions that already exist. For this reason we discuss regarding our needs.

1.4.1 Collaborating Network

We assume there is a network of computers which are available to run the tasks, each node is running an instance of the application. We will propose our collaboration and data transfer algorithm between them later.

1.4.2 Data Characteristics

We need to discuss more about the data. In our scientific context data is mostly numerical and explains characteristics of physical particles such as atoms and molecules. These data is being used to simulate collections of particles called models. Although our work is not dependent on these, they help us to understand the the definition of the data that we often refer to in this report. One important aspect of the data that we are interested in is that it is not critical and we can reproduce it.

1.4.3 Data Transfer

We assume a data transfer approach is already in place. This could be any file system which supports network storage. Rather than going into details of how data could be transferred more efficiently, we will focus on finding which data to be transferred and from which computer to which destination.

1.4.4 Workflow

In contrast to data we are interested in workflow. We want to find a reliable approach to access and update state of our workflow on any arbitrary node

which is part of our collaborative network.

Approaches

During this work we consider three different approaches toward preparing required data for operations.

Conventional Approach in this approach we put the required data on a network file system and all application instances will access it there. We will utilize an NFS mounted file system.

Centralized Approach in this approach we will have a central instance which will orchestrate operation delegation and operation output forwarding to other nodes.

Decentralized Approach in this approach we will eliminate the orchestrator node and the network of application instances should collaborate in a decentralized fashion to keep track of data and control flow for each task.

For every approach we will run performance tests and we will compare the results.

Method

We will discuss scenarios in [4](#). For each scenario we will analyze the possible combinations of data and operations and we will discuss how to deliver the input data and where to store output data. We will discuss workflow management in chapter [5](#) and data transfer in [6](#).

Chapter 2

Rough Ideas

This chapter contains very raw ideas to address main requirements i.e. distributed workflow management and intelligent data transfer.

2.1 Intelligent Data Transfer: Use Case One

By *Intelligent Data Transfer* we mean an approach that minimizes required data transfers between application¹ instances.

In the most basic use case² we run a script³ which consists of two linear operations. Each operation consumes data and generates data. A third operation needs both generated data two operate on and generate the third and final data.

The script is data driven. It means that it contains a number of steps and for every step it needs appropriate data to run the desired operations⁴. We assume that the script will run on *Node 1* and required data *DataSet1* and *DataSet2* are located on *Node 2* and *Node 3* respectively. Therefor *Node 1* have to initiate operations on the other machines.

2.1.1 Identical Instances

We assume that on each machine of the network the same instance of our imaginary program is running which is capable of running all operations including A, B and C. The only consideration is the availability of DatSets, they are not available on all machines.

¹To be defined

²To be added later and referenced here

³To be defined and added to the terminology, terminology itself has to be defined

⁴To be defined

A linear operation (not clear to my self how to write it):

$$Operation(A, B) = Operation(A) + Operation(B)$$

$$DataSet^A = Operation^A(DataSet^1)$$

$$DataSet^B = Operation^B(DataSet^2)$$

$$DataSet^C = Operation^C(DataSet^A, DataSet^B)$$

Assuming that operations A and B will run on the machines which contain the required data, a number of questions arise here:

1. On which machine operations C should run? A, B or C?
2. On How to transfer the required data to that machine in an optimized way?

2.1.2 The Idea

First of all we assume that we have the information about the DataSets available on all of the machines i.e. in form of a distributed table with entries containing the node address and DataSet id. Based on this information the application can decide if it has the required data or not.

Based on this algorithm (to be defined) the initial application delegates operations to the other nodes (instances of the same program), where the data is available. Our distributed workflow manager (to be defined) will synchronize the information on these running operation and will label the output data and will add it to the distributed data table.

After finishing operations A and B we will run operation C in either of these nodes, because the required data is partially available on these nodes. Then we have to transfer the rest of the data to one of these nodes to run the operation C which needs both parts simultaneously.

Using Prior Art At this point we can take advantage of existing Distributed File Systems (DFS) to make the data available for operation C. We can then eliminate the complexity of data transfer between these two nodes and delegate it to existing distributed file systems. The main point is we don't rely on DFS for all of our decision making part but we explicitly make the decision to run operation A and B on specific nodes and then for the last part we use a meta disk or universal disk concept to deliver the remaining data for operation C.

Chapter 3

Literature

In this chapter we will go through a number of existing solutions and we will discuss their efficiency and deployment complexity. Before that we introduce the parameters which are important for us. Then we will assess each solution against the introduced parameter set.

3.1 Parameters

There are many factors that we need to take into account before introducing our parameters. To name a few:

- Data transfer cost
- Data reproduction cost
- Type of operation on data, linear or non-linear
- Replication
- Deployment complexity
- Fault tolerance
- Portability
- Performance according to number of distributed nodes and size of files

3.2 Data Storage Systems

SAME POINTS SHOULD BE DISCUSSED FOR EACH APPROUACH (FROM OUR POINT OF VIEW) e.g. EFFICIENCY, COMPLEXITY, DISTRIBUTED APPLICATION ACCESS, APPLICATION AWARENESS, POSSIBLE DATA ACCESS SCENARIOS, SCALIBILITY, DATA TRANSFER, FAULT TOLERANCE, ACCESS CONTROL

3.3 Distributed File Systems

One way to achieve fault tolerant and reliable data storage and access is to use distributed file systems (DFS). In this case the data will be replicated over a network of storage servers with different magnitudes based on the underlying file system. We will discuss a number of free and open source solutions.

3.3.1 Hadoop Distributed File System (HDFS)

“The Hadoop Distributed File System (HDFS) is a distributed file system designed to run on commodity hardware.” [7, tp. 3]

“Hadoop1 provides a distributed filesystem and a framework for the analysis and transformation of very large data sets using the MapReduce [1] paradigm.” [6]

“HDFS stores metadata on a dedicated server, called the NameNode. Application data are stored on other servers called DataNodes.” [6]

Deployment Complexity src:<http://hadoop.apache.org/docs/r0.18.3/quickstart.html> needs Java 1.5.x ssh and sshd and rsync. Three basic modes are available: Local, Pseudo-Distributed and Fully Distributed mode. XML configuration, installation of Local and Pseudo Distributed modes are almost straight forward, for fully distributed note extra steps are required (official doc link is dead).

Efficiency

Fault Tolerance “Hardware failure is the norm rather than the exception.” “Each DataNode sends a Heartbeat message to the NameNode periodically.” “The DataNodes in HDFS do not rely on data protection mechanisms such as RAID to make the data durable. Instead, like GFS, the file content is replicated on multiple DataNodes for reliability.” [6]

Portability “HDFS has been designed to be easily portable from one platform to another.”

Robustness “The primary objective of HDFS is to store data reliably even in the presence of failures.”

Accessibility

1. FS Shell
2. DFSAdmin
3. Browser

Applicability There is a good document here: http://hadoop.apache.org/docs/r0.18.0/hdfs_design.pdf Hints: HADOOP is for big data and the programming should be different (map/reduce) and it does not look suitable for our use cases and requirements. The burden would be so high that we will waste a lot of resources. I have to put these in scientific words with more logic and references to sizes that we need and more numbers.

Users have to program their applications using Java and Hadoop to take advantage of distributed computing features in Hadoop MapReduce and HDFS. Cites? Hadoop website? <https://infosys.uni-saarland.de/publications/BigDataTutorial.pdf>

3.3.2 XTREEMFS

Chapter 4

Scenarios

There are a number of possible use cases in our problem domain. To demonstrate these cases we assume we have a number of nodes and datasets respectively, but they are not necessarily on the same nodes. In the following paragraphs we explain possible combinations of operations, nodes and datasets.

4.1 Operation

In every scenario we want to run an operation which could be linear or non-linear.

4.1.1 Linear Operation

Being linear means that the operation could be broken into its components and then run in parallel or series. Here is algebraic notation of a linear operation which acts on two datasets:

$$Operation(A + B) = Operation(A) + Operation(B)$$

Being linear or non-linear only matters when we have to operate on more than one dataset.

4.1.2 Non-linear Operation

In contrast to linear there is non-linear operation. This means that this kind of operation has dependant parts and those parts could not run in parallel:

$$Operation(A + B) \neq Operation(A) + Operation(B)$$

4.2 Datasets

For each operation we need one or more datasets which may be available on the same node that wants to run the operation initially or could reside on other nodes.

4.2.1 Input

Input files are normally not mission critical and could be reproduced.

4.2.2 Output

Operations create output datasets which normally are small in size, therefore we ignore the transfer cost of operation results in our work.

4.2.3 Data Locationing

We consider three different approaches toward preparing required data for operations.

Conventional Approach

in this approach we put the required data on a network file system and all application instances will access it there. We will utilize an NFS mounted file system.

Centralized Approach

in this approach we will have a central instance which will orchestrate operation delegation and operation output forwarding to other nodes.

Decentralized Approach

in this approach we will eliminate the orchestrator node and the network of application instances should collaborate in a decentralized fashion to keep track of data and control flow for each task.

For every approach we will run performance tests and we will compare the results.

Method

We will discuss scenarios in 4. For each scenario we will analyze the possible combinations of data and operations and we will discuss how to deliver the input data and where to store output data. We will discuss workflow management in chapter 5 and data transfer in 6.

4.3 Decision Making

The main decision that we need to make at every scenario is whether we should transfer the required data or we need to delegate the operation to an instance on a node which already has the data. To make a decision we need to answer a number of questions. First we need to know the location of the data:

1. Is the data available locally?
2. If not, is the data available on another node? – Here only the physical location of data matters not the instance controlling it.

4.3.1 Scenario 1

In this scenario we have a linear operation, e.g. Op^A on $Node^A$ which requires $Dataset^1$.

Conditions

$Dataset^1$ is not available on $Node^A$.

Consequences

With these conditions we either should transfer $Dataset^1$ to local node or in case of availability delegate Op^A to the node which already has $Dataset^1$.

Chapter 5

Workflow Management

Chapter 6

Data Transfer

References

- [1] Jeffrey Dean and Sanjay Ghemawat. *MapReduce: Simplified Data Processing on Large Clusters*. Proc. Sixth Symposium on Operating System Design and Implementation, 2004.
- [2] Jeffrey Dean and Sanjay Ghemawat. “MapReduce: Simplified Data Processing on Large Clusters”. In: *Commun. ACM* 51.1 (Jan. 2008), pp. 107–113. ISSN: 0001-0782. DOI: [10.1145/1327452.1327492](https://doi.org/10.1145/1327452.1327492). URL: <http://doi.acm.org/10.1145/1327452.1327492>.
- [3] C. Moretti et al. “All-pairs: An abstraction for data-intensive cloud computing”. In: *Parallel and Distributed Processing, 2008. IPDPS 2008. IEEE International Symposium on*. Apr. 2008, pp. 1–11. DOI: [10.1109/IPDPS.2008.4536311](https://doi.org/10.1109/IPDPS.2008.4536311).
- [4] S. Pandey et al. “A Particle Swarm Optimization-Based Heuristic for Scheduling Workflow Applications in Cloud Computing Environments”. In: *Advanced Information Networking and Applications (AINA), 2010 24th IEEE International Conference on*. Apr. 2010, pp. 400–407. DOI: [10.1109/AINA.2010.31](https://doi.org/10.1109/AINA.2010.31).
- [5] K. Ranganathan and I. Foster. “Decoupling computation and data scheduling in distributed data-intensive applications”. In: *High Performance Distributed Computing, 2002. HPDC-11 2002. Proceedings. 11th IEEE International Symposium on*. 2002, pp. 352–358. DOI: [10.1109/HPDC.2002.1029935](https://doi.org/10.1109/HPDC.2002.1029935).
- [6] *The Hadoop Distributed File System*. URL: <http://www.aosabook.org/en/hdfs.html>.
- [7] *The Hadoop Distributed File System: Architecture and Design*. 2007. URL: http://hadoop.apache.org/docs/r0.18.0/hdfs_design.pdf.